

```
In [1]: import numpy as np
        from scipy import stats
        import matplotlib.pyplot as plt
        import pystan
```

1. Call Center Data

Data: waiting times for the 13th hour of a day in a call center

Prior distribution: Gamma distribution with $\alpha = 1$ and $\beta = 0.25$

Likelihood function: exponential with parameter λ

Parameters: rate λ

Posterior: Gamma distribution over λ

```
In [2]: ## import the dataset (code from call_center_solution.ipynb)
        waiting_times_day = np.loadtxt('call_center.csv')

        # Split the data into 24 separate series, one for each hour of the day
        current_time = 0
        waiting_times_per_hour = [[] for _ in range(24)] # Make 24 empty lists, one per hour
        for t in waiting_times_day:
            current_hour = int(current_time // 60)
            current_time += t
            waiting_times_per_hour[current_hour].append(t)

        # use just the 13th hour of the day
        waiting_times_hour = waiting_times_per_hour[13]
```

```
In [3]: # define the data for the 13th hour of the day
        call_center_data = {
            '13': {
                'alpha': 1, # fixed prior hyperparameters for the
                'beta': 0.25, # gamma distribution
                'num_calls': len(waiting_times_hour), # number of calls coming in
                'waiting_times': waiting_times_hour} # data set on waiting times
        }
```

```
In [4]: # define stan code
calls_stan_code = """

// The data block contains all known quantities - typically the observed
// data and any constant hyperparameters.
data {
  int<lower=1> num_calls; // number of calls
  real<lower=0> waiting_times[num_calls]; // waiting times
  real<lower=0> alpha; // fixed prior hyperparameter
  real<lower=0> beta; // fixed prior hyperparameter
}

// All unknown quantities, in this case the waiting time lambda
parameters {
  real lambd; // rate lambda for the exponential
}

// The model block contains all probability distributions in the model.
model {
  lambd ~ gamma(alpha, beta); // prior over lambda
  for(i in 1:num_calls) {
    waiting_times[i] ~ exponential(lambd); // likelihood function
  }
}

"""
```

```
In [5]: # define stan model
calls_stan_model = pystan.StanModel(model_code=calls_stan_code)
```

INFO:pystan:COMPILING THE C++ CODE FOR MODEL anon_model_50fe82232f7cd8b2736c8b3bf1959587 NOW.

```
In [6]: # evaluate model with the dataset and print parameter values
calls_stan_results = calls_stan_model.sampling(data=call_center_data['13'])
print(calls_stan_results.stansummary(pars=['lambda'], probs=[0.01, 0.5, 0.99]))
```


Inference for Stan model: anon_model_50fe82232f7cd8b2736c8b3bf1959587.
 4 chains, each with iter=2000; warmup=1000; thin=1;
 post-warmup draws per chain=1000, total post-warmup draws=4000.

	mean	se_mean	sd	1%	50%	99%	n_eff	Rhat
lambda	6.35	8.0e-3	0.32	5.61	6.34	7.15	1618	1.0

Samples were drawn using NUTS at Fri Oct 16 00:42:46 2020.
 For each parameter, n_eff is a crude measure of effective sample size,
 and Rhat is the potential scale reduction factor on split chains (at
 convergence, Rhat=1).

```
In [7]: # plot a histogram of the distribution of lambdas from the model
posterior_calls_samples = calls_stan_results.extract()
plt.hist(posterior_calls_samples['lambda'], bins=50, density=True)
plt.xlabel("lambda")
plt.title('Sampled posterior probability density for lambda')
print(
    "Posterior 98% confidence interval for lambda:",
    np.percentile(posterior_calls_samples['lambda'], [1, 99]))
plt.show()
```

Posterior 98% confidence interval for lambda: [5.61525608 7.14944725]



2. Normal likelihood with normal-inverse-gamma prior

Prior distribution: normal inverse gamma distribution

Likelihood function: normal with parameters μ and σ^2

Parameters: mean μ and variance σ^2

Posterior: normal inverse gamma

```
In [8]: #define dataset
raw_norm_inv_gamma_data = np.array([3.54551763569501, 4.23799861761927, 4.72138425951628, -0.69226532036
norm_inv_gamma_data = {
    'mu': 0, # prior mean centered at 0
    'nu': 0.054, # nu indicates the uncertainty of the prior mean
    'alpha' : 1.12, # alpha and beta govern the marginal prior over the variance
    'beta' : 0.4,
    'data_length': len(raw_norm_inv_gamma_data), # number of data points
    'norm_data': raw_norm_inv_gamma_data} # data set
```

```
In [9]: # define the stan model
norm_inv_gamma_stan_code = """

// The data block contains all known quantities - typically the observed
// data and any constant hyperparameters.
data {
    int<lower=1> data_length; // number of data points
    real norm_data[data_length]; // data points
    real mu; // fixed prior hyperparameter
    real nu; // fixed prior hyperparameter
    real alpha; // fixed prior hyperparameter
    real beta; // fixed prior hyperparameter
}

// All unknown quantities, in this case the mean and standard deviation of the data
parameters {
    real x; // mean of the data
    real<lower=0> sigma2; // variance of the data
}

// The model block contains all probability distributions in the model.
model {
    sigma2 ~ inv_gamma(alpha, beta); // prior over mean
    x ~ normal(mu,sqrt(sigma2/nu)); //prior over variance
    for(i in 1:data_length) {
        norm_data[i] ~ normal(x,sqrt(sigma2)); // likelihood function
    }
}

"""
```

In [10]: *# evaluate stan model*

```
norm_inv_gamma_stan_model = pystan.StanModel(model_code=norm_inv_gamma_stan_code)
```

INFO:pystan:COMPILING THE C++ CODE FOR MODEL anon_model_775fbb51b12e7e2ff371e277ab8fbfec NOW.

In [11]: *# input data to stan model and print the results*

```
norm_inv_gamma_stan_results = norm_inv_gamma_stan_model.sampling(data=norm_inv_gamma_data)
print(norm_inv_gamma_stan_results.stansummary(pars=['x', 'sigma2'], probs=[0.025, 0.5, 0.975]))
```

Inference for Stan model: anon_model_775fbb51b12e7e2ff371e277ab8fbfec.

4 chains, each with iter=2000; warmup=1000; thin=1;

post-warmup draws per chain=1000, total post-warmup draws=4000.


	mean	se_mean	sd	2.5%	50%	97.5%	n_eff	Rhat
x	3.06	2.4e-3	0.14	2.8	3.06	3.33	3163	1.0
sigma2	3.62	6.7e-3	0.36	2.99	3.59	4.36	2851	1.0

Samples were drawn using NUTS at Fri Oct 16 00:43:47 2020.

For each parameter, n_eff is a crude measure of effective sample size, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat=1).

```
In [12]: # extract results from stan and get 10 random samples
posterior_norm_inv_gamma_samples = norm_inv_gamma_stan_results.extract()
samples_i = np.random.randint(4000,size=10)
xs = posterior_norm_inv_gamma_samples['x'][samples_i]
sigma2s = posterior_norm_inv_gamma_samples['sigma2'][samples_i]

# plot the normal distributions corresponding to the samples
plt.figure(figsize=(12, 6))
plot_x = np.linspace(-15, 15, 500)
for i in range(len(xs)):
    plot_y = stats.norm.pdf(plot_x, loc=xs[i], scale=np.sqrt(sigma2s[i]))
    plt.plot(plot_x, plot_y)
plt.title('%i samples from a normal-inverse-gamma posterior distribution' % len(xs))
plt.show()
```



3. Log-normal HRTEM data

Data: particle sizes in nanometers

Prior distribution: normal inverse gamma distribution

Likelihood function: normal with parameters μ and σ^2


Parameters: mean μ and variance σ^2

Posterior: normal inverse gamma

```
In [13]: raw_hrtem = np.loadtxt('hrtem.csv')
log_hrtem = np.log(raw_hrtem)
hrtem_data = {
    'mu': 2.3,
    'nu': 0.1,
    'alpha': 2,
    'beta': 5,
    'data_length': len(log_hrtem), # number of data points
    'hrtem_data': log_hrtem} # data set
```



```
In [14]: plt.hist(log_hrtem)  
plt.show()
```



```
In [15]: hrtem_stan_code = """

// The data block contains all known quantities - typically the observed
// data and any constant hyperparameters.
data {
  int<lower=1> data_length; // number of data points
  real hrtem_data[data_length]; // data points
  real mu; // fixed prior hyperparameter
  real nu; // fixed prior hyperparameter
  real<lower=0> alpha; // fixed prior hyperparameter
  real<lower=0> beta; // fixed prior hyperparameter
}

// All unknown quantities, in this case the mean and standard deviation of the data
parameters {
  real x; // mean of the data
  real<lower=0> sigma2; // variance of the data
}

// The model block contains all probability distributions in the model.
model {
  sigma2 ~ inv_gamma(alpha, beta); // prior over mean
  x ~ normal(mu, sqrt(sigma2/nu)); //prior over variance
  for(i in 1:data_length) {
    hrtem_data[i] ~ normal(x, sqrt(sigma2)); // likelihood function
  }
}

"""
```

```
In [16]: hrtem_stan_model = pystan.StanModel(model_code=hrtem_stan_code)
```

INFO:pystan:COMPILING THE C++ CODE FOR MODEL anon_model_0c47c1ab3212c3dd8140560a3a5e5e9f NOW.

```
In [17]: hrtem_stan_results = hrtem_stan_model.sampling(data=hrtem_data)
print(hrtem_stan_results.stansummary(pars=['x', 'sigma2'], probs=[0.025, 0.5, 0.975]))
```

Inference for Stan model: anon_model_0c47c1ab3212c3dd8140560a3a5e5e9f.
 4 chains, each with iter=2000; warmup=1000; thin=1;
 post-warmup draws per chain=1000, total post-warmup draws=4000.

	mean	se_mean	sd	2.5%	50%	97.5%	n_eff	Rhat
x	1.89	5.4e-4	0.03	1.83	1.89	1.96	3453	1.0
sigma2	0.5	5.4e-4	0.03	0.44	0.5	0.56	3407	1.0

Samples were drawn using NUTS at Fri Oct 16 00:44:58 2020.
 For each parameter, n_eff is a crude measure of effective sample size,
 and Rhat is the potential scale reduction factor on split chains (at
 convergence, Rhat=1).

```

In [18]: # extract 10 samples from the distribution
posterior_hrtem_gamma_samples = hrtem_stan_results.extract()
hrtem_samples_i = np.random.randint(4000,size=10)
hrtem_xs = posterior_hrtem_gamma_samples['x'][hrtem_samples_i]
hrtem_sigma2s = posterior_hrtem_gamma_samples['sigma2'][hrtem_samples_i]

# plot the log-normal together with the data
plt.figure(figsize=(12, 6))
plot_x = np.linspace(0, 30, 500)
plt.hist(raw_hrtem, bins=20, density=True, alpha =0.5)
for i in range(len(xs)):
    plot_y = stats.lognorm.pdf(plot_x, np.sqrt(hrtem_sigma2s[i]), scale= np.exp(hrtem_xs[i]))
    plt.plot(plot_x, plot_y)
plt.title('%i samples of the posterior log-normal pdf' % len(hrtem_xs))
plt.show()

```

